

Models of Computation for Homomorphic Encryption. TFHE - Chimera

Nicolas Gama and Mariya Georgieva Belorgey



Based on joint work with: C. Boura, I. Chillotti, M. Izabachene, D. Jetchev

Little history of FHE

Constant-time Privacy-preserving computations

- FHE programs behave essentially like circuits.
- 2009: Bootstrapping (evaluate one homomorphic NAND)
- The overhead between plaintext circuit and FHE is $O(1)$ in time and memory

42: That's it!! CQFD?

- 2009: First bootstrapping = hours of computations
- 2009-now: Just wait for scientific progress... and brace for impact!!

Little history of FHE


Constant-time Privacy-preserving computations

- FHE programs behave essentially like circuits.
- 2009: Bootstrapping (evaluate one homomorphic NAND)
- The overhead between plaintext circuit and FHE is $O(1)$ in time and memory


42: That's it!! CQFD?

- 2009: First bootstrapping = hours of computations
- 2009-now: Just wait for scientific progress... and brace for impact!!


Bootstrapping: the beginnings



Bootstrapping: the beginnings



Bootstrapping: the beginnings




Properties of a gate-bootstrapping ciphertext

Properties of a gate-bootstrapping ciphertext


- a message = one bit

Composition rules: gates


- constant gates (0,1)
- unary gates (copy, not)
- binary gates (and, or, nand, nor, ...)
- selector gate (mux)




Little history: are we there yet?




Little history: are we there yet?



Little history: are we there yet?



Little history: are we there yet?



Gate bootstrapping: pros/cons

Pros of gate bootstrapping

- circuits are standard and easy to generate
- asymptotically $O(1)$ optimal w.r.t. plaintext circuit

Cons of gate bootstrapping

- the $O(1)$ theoretical overhead factor is huge in practice
 - **timing** 10ms vs. 1 nanosec per cycle (vs. 1 picosec physical),
 - **size** 20kB ciphertext per bit.
 - **parallelization** 100 gates in parallel vs. billion gates in parallel.

The reality: 2023

- Only use-cases that take a fraction of second in plaintext are feasible via only Gate Bootstrapping.
- Practical FHE requires a plan B!

Gate bootstrapping: pros/cons

Pros of gate bootstrapping

- circuits are standard and easy to generate
- asymptotically $O(1)$ optimal w.r.t. plaintext circuit

Cons of gate bootstrapping

- the $O(1)$ theoretical overhead factor is huge in practice
 - **timing** 10ms vs. 1 nanosec per cycle (vs. 1 picosec physical),
 - **size** 20kB ciphertext per bit.
 - **parallelization** 100 gates in parallel vs. billion gates in parallel.

The reality: 2023

- Only use-cases that take a fraction of second in plaintext are feasible via only Gate Bootstrapping.
- Practical FHE requires a plan B!

Plan B: LHE to the rescue of FHE

FHE Fully Homomorphic Encryption ($O(1)$ from optimal)

\exists crypto params s.t. \forall circuit C , we can evaluate C homomorphically.

LHE Leveled Homomorphic Encryption (not $O(1)$ from optimal)

\forall circuit C , \exists crypto params s.t. we can evaluate C homomorphically.

Chimera: Using FHE to boost LHE

- FHE and LHE are not mutually exclusive, they should be used together!
- Some LHE schemes are much faster for massive low multiplication-depth arithmetic in practice (integer (BFV) or FP (CKKS)).
- Other LHE schemes are quite good for evaluating automata (RGSW, TFHE).
- Bootstrapping is quite good at evaluating LUTs and univariate non-linear functions, like conversions.

Compilation

Nice to have


Let the user write the desired program as a high-level pseudocode.

```
def myfunction(x, y, z, bigvector, bigmatrix)
  a := 2 * x + 3 * y * z mod 15
  c := 30 * cos(a / 8.)
  return c * bigmatrix * bigvector
```

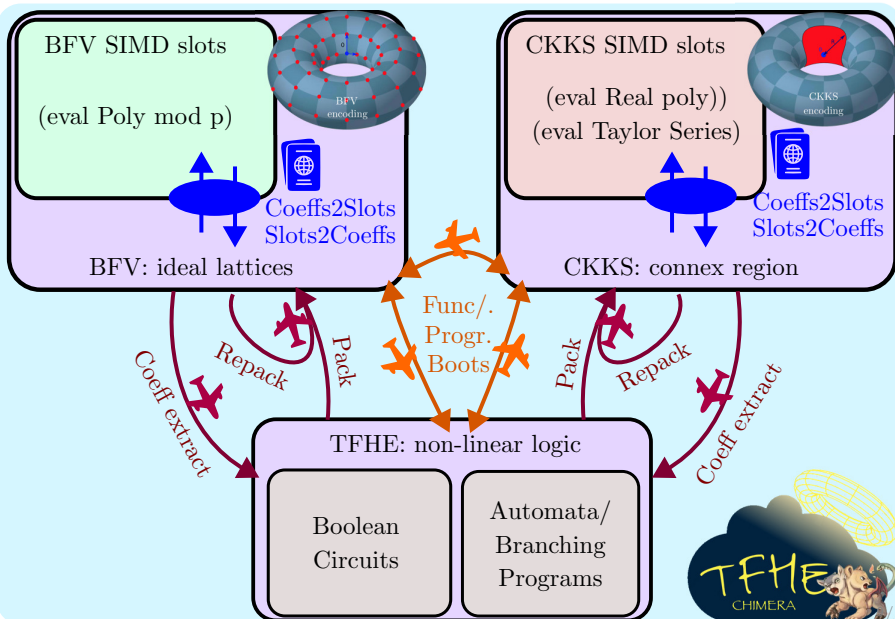
Nice to have

Let the user write the desired program as a high-level pseudocode.

```
def myfunction(x, y, z, bigvector, bigmatrix)
  a := 2 * x + 3 * y * z mod 15
  c := 30 * cos(a / 8.)
  return c * bigmatrix * bigvector
```



Chimera World Map



Properties of a leveled ciphertext

Properties of a leveled HE ciphertext

- a message (encoded in a polynomial)
 - one bit? one integer?
 - a vector of integers mod p ?
 - a vector of floats?
- a homomorphic budget: 3 equivalent definitions
 - noise rate: $0 < \alpha < 1$
 - homomorphic budget: $-\log_2(\alpha) \geq 0$ points
 - homomorphic level: 1 level ≈ 30 points

Homomorphic budget (a.k.a. level, noise-rate)

Noise rate and homomorphic budget

A noise rate $a < 1$ corresponds to a Homomorphic budget of $-\log_2(\alpha) > 0$ points, and quantifies the number of homomorphic operations that can be carried out on a FHE ciphertext.

	efficiency	HE operations
$-\log_2(\alpha) \approx 0$	small ciphertext, small key, fast operations	exhausted
$-\log_2(\alpha) \approx 30$	native 32-bit arithmetic	1 multiplications
$-\log_2(\alpha) \approx 60$	native 64-bit arithmetic	2 multiplications
$-\log_2(\alpha) \approx 300$	slow 300-bit arithmetic	10 multiplications

Rule of thumb

- more points = more homomorphic power. But bigger ciphertexts, larger keys, larger arithmetic, slower operations.
- decreasing the homomorphic budget points is easy ("*modulus switching/rescaling*").
- increasing it is much harder ("*bootstrapping*").

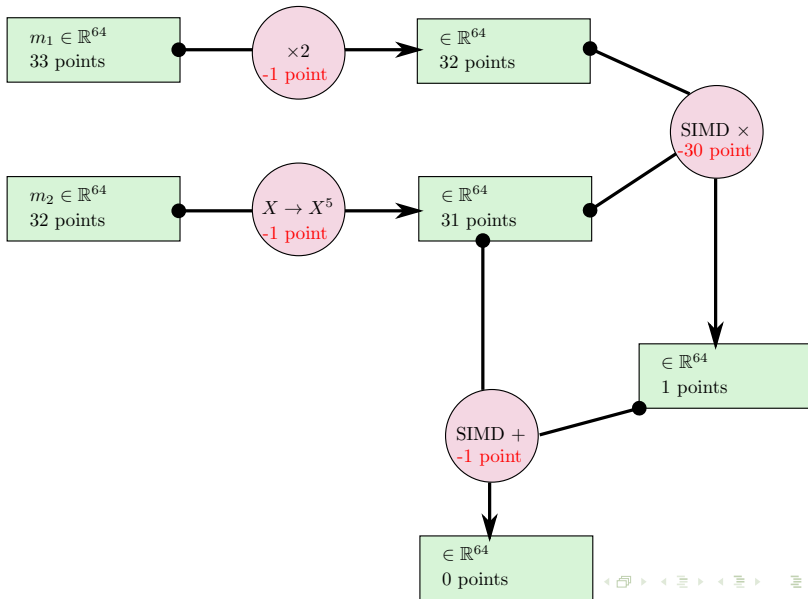
Operations and composition rules

Arithmetic circuit model of computation

A graph of polynomial arithmetic operations. Each operation impact the noise (so the homomorphic budget):

Operation	ciphertext type	homomorphic budget impact
slot-wise product:	RLWE	-30 points
sum	RLWE	-1 point
public rotation:	RLWE	-0 point
linear combinations $\sum e_i c_i$:	RLWE	$-\log(1 + \ e\ _1)$ points
substitution with X^t :	RLWE	-1 point

Operations and composition rules




Bootstrapping rules


Arithmetic Circuit: bootstrapping

- **Input:** a ciphertext with (nearly) depleted budget
- **Output:** a ciphertext of the same message with much larger budget
- **Restrictions/Rules:**
 - Message space, maximal FP precision, modulus.
 - Minimum input level/points, output level/points
 - Running time


Bootstrapping: example in the literature



Bootstrapping is per type of message space: real, mod, binary.
1 Bootstrapping proposes a fixed input/output level



Bootstrapping may use all the slots, or just a fewer number.
Bootstrapping may also require a minimum input level > 0



Arithmetic circuits: Pros and Cons

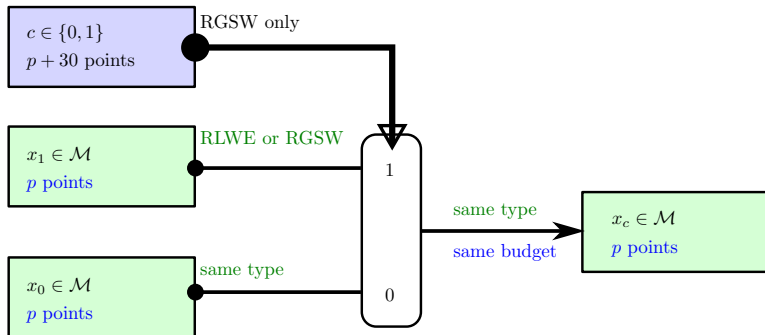
Arithmetic circuits: Pros

- the "*assembly language*" of LWE: no loss, no overhead
- good for executing SIMD arithmetic use-cases
- Newest bootstrapping have usually a fast amortized time per slot.

Arithmetic circuits: Cons

- General use-cases are hard to convert to polynomial arithmetic circuits. (think CPU vs. GPU)
- The rare use-cases that work are already described "in assembly"

RGSW-based private selector circuits



data Input and output have the same homomorphic budget!

LUT evaluation

LookUp Tables (LUT) to evaluate arbitrary functions:


$$f: \mathbb{B}^d \longrightarrow \mathbb{T}^s$$

$$x = (x_0, \dots, x_{d-1}) \longmapsto f(x) = (f_0(x), \dots, f_{s-1}(x))$$

Example with $d = 3$ and $s = 2$

x_0	x_1	x_2	f_0	f_1
0	0	0	0.5	0.3
1	0	0	0.25	0.7
0	1	0	0.1	0.61
1	1	0	0.83	0.9
0	0	1	0.23	0.47
1	0	1	0.67	0.42
0	1	1	0.78	0.12
1	1	1	0.35	0.95

Evaluation via MUX tree



LUT evaluation

LookUp Tables (LUT) to evaluate arbitrary functions:


$$f: \mathbb{B}^d \longrightarrow \mathbb{T}^s$$

$$x = (x_0, \dots, x_{d-1}) \mapsto f(x) = (f_0(x), \dots, f_{s-1}(x))$$

Example with $d = 3$ and $s = 2$

x_0	x_1	x_2	f_0	f_1
0	0	0	0.5	0.3
1	0	0	0.25	0.7
0	1	0	0.1	0.61
1	1	0	0.83	0.9
0	0	1	0.23	0.47
1	0	1	0.67	0.42
0	1	1	0.78	0.12
1	1	1	0.35	0.95

Evaluation via MUX tree



LUT evaluation

LookUp Tables (LUT) to evaluate arbitrary functions:


$$f: \mathbb{B}^d \longrightarrow \mathbb{T}^s$$

$$x = (x_0, \dots, x_{d-1}) \longmapsto f(x) = (f_0(x), \dots, f_{s-1}(x))$$

Example with $d = 3$ and $s = 2$

x_0	x_1	x_2	f_0	f_1
0	0	0	0.5	0.3
1	0	0	0.25	0.7
0	1	0	0.1	0.61
1	1	0	0.83	0.9
0	0	1	0.23	0.47
1	0	1	0.67	0.42
0	1	1	0.78	0.12
1	1	1	0.35	0.95

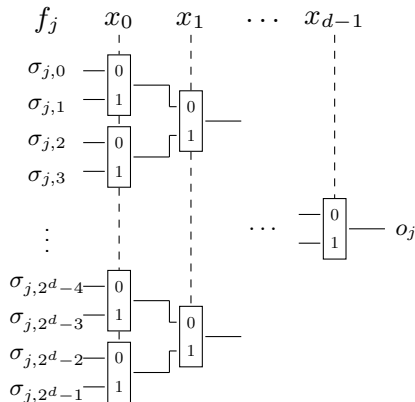
Evaluation via MUX tree



LUT evaluation

How to evaluate it?

x_0	\dots	x_{d-1}	f_0	\dots	f_{s-1}
0	\dots	0	$\sigma_{0,0}$	\dots	$\sigma_{s-1,0}$
1	\dots	0	$\sigma_{0,1}$	\dots	$\sigma_{s-1,1}$
0	\dots	0	$\sigma_{0,2}$	\dots	$\sigma_{s-1,2}$
1	\dots	0	$\sigma_{0,3}$	\dots	$\sigma_{s-1,3}$
\vdots	\dots	\vdots	\vdots	\vdots	\vdots
0	\dots	1	$\sigma_{0,2^d-4}$	\dots	$\sigma_{s-1,2^d-4}$
1	\dots	1	$\sigma_{0,2^d-3}$	\dots	$\sigma_{s-1,2^d-3}$
0	\dots	1	$\sigma_{0,2^d-2}$	\dots	$\sigma_{s-1,2^d-2}$
1	\dots	1	$\sigma_{0,2^d-1}$	\dots	$\sigma_{s-1,2^d-1}$



DFA versus det-WFA

TFHE 2016: DFA (deterministic finite automata)

- Decisional: returns **accepted** (1) or **rejected** (0)


TFHE 2017: det-WFA (deterministic weighted finite automata)

- Computational: returns a **weight** in $\mathbb{T}_N[X]$

Weights act like a "memory" that stores the result all along the evaluation

DFA versus WFA

Deterministic Finite Automata (DFA)




Acceptance

"00101" → False

"10111" → True

Deterministic Weighted Finite Automata (det-WFA)




Weight Computation

"00101" → (2,0)

"10111" → (4,1)


DFA computational models

mirror(\mathcal{L})
rev. det. autom.





DFA computational models

mirror(\mathcal{L})
rev. det. autom.





DFA computational models

mirror(\mathcal{L})
rev. det. autom.





DFA computational models

mirror(\mathcal{L})
rev. det. autom.



DFA computational models

mirror(\mathcal{L})
rev. det. autom.



Computation of the maximum

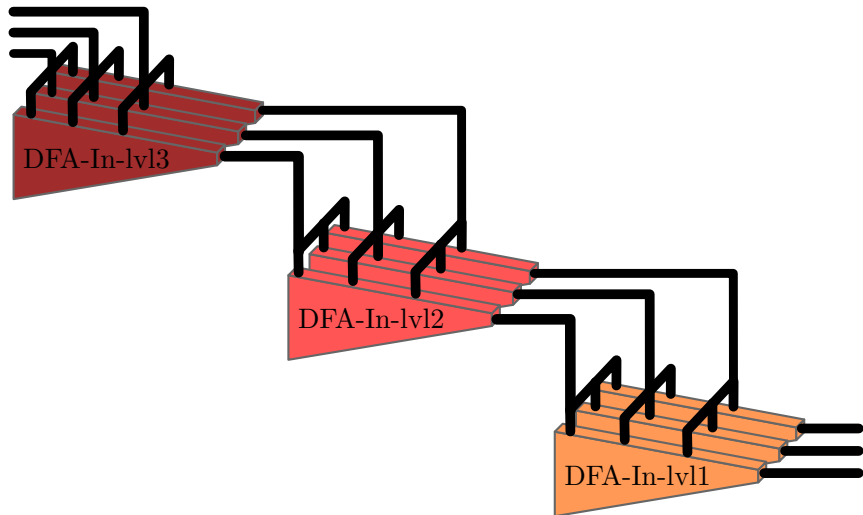
Example: evaluation of $m = MAX(x, y)$

Let $x = (x_1, \dots, x_n)$ and $y = (y_1, \dots, y_n)$.

We want to compute $m = (m_1, \dots, m_n) = MAX(x, y)$.

- **DFA**: evaluate n **DFA**, one per output bit
- **Det-WFA**: evaluate **1 det-WFA**, the result given in a single path


Arbitrary long Composition of automata?



TFHE in Circuit Bootstrap mode

Circuit bootstrapping CGGI2017

- Take advantage that the message space is binary
- And that input/output levels are very low ($0 \rightarrow 60$ points)
- Reconstruct a TRGSW encryption directly from its internal structure [CGGI17] rather than as the output of larger homomorphic operations (see [GSW13], [AP14] constructions).



Accepted inputs:


A TLWE ciphertext on binary message space $\{0, \frac{1}{2}\}$

One coefficient of a TRLWE ciphertext over $\{0, \frac{1}{2}\}^N$

One coefficient of a TRGSW ciphertext over $\{0, 1\}^N$


Circuit mode versus Gate Bootstrap mode

Gate bootstrapping mode



versus (or together with?)

Circuit bootstrapping mode



In Summary: The Chimera VM

Logical unit

- Digital circuits
- Lookup Tables
- Deterministic Automata (finite and weighted)

Heavy arithmetic unit

- SIMD fixed-point and modular unit
- Support also convolution, big-integers

Composability, Compilation

- A rich VM capturing all the capabilities of RLWE-based FHE.
- Immediate link to the lattice geometry and its security.
- Is it feasible to compile for this programming model?


Anatomy of a ciphertext

WHAT ARE LEGO BRICKS MADE OF?

LEGO BRICKS



Up until 1963, cellulose acetate was used to make Lego bricks and parts. Lego bricks are now made from acrylonitrile butadiene styrene (ABS). ABS is less subject to warping and colour fading.



HOW LEGO IS MADE



ABS GRANULES ADDED

Macrollex dyes are added to ABS to colour it.

GRANULES HEATED TO 230 °C (450 °F)

MELTED PLASTIC FED INTO MOLDS

In 2014 more than 60 billion Lego pieces were made.

OTHER LEGO PARTS



ABS is opaque, so a polycarbonate polymer has to be used for transparent Lego parts. For leaves, bushes and trees, Lego has recently started using polyethylene derived from sugar cane.



POLYETHENE

SBS

Tyres and elastic materials are made from styrene butadiene styrene (SBS).



© Andy Brunning/Compound Interest 2018 - www.compoundchem.com | Twitter: @compoundchem | FB: www.facebook.com/compoundchem
 This graphic is shared under a Creative Commons Attribution-NonCommercial-NoDerivatives licence.



Integer/Real/Complex Polynomials

Ring of polynomials with coefficients $\in \mathbb{Z}, \mathbb{R}$ or $\mathbb{C} \bmod X^N + 1$:

$$\mathbb{Z}_N[X] = \mathbb{Z}[X]/(X^N + 1)$$

$$\mathbb{R}_N[X] = \mathbb{R}[X]/(X^N + 1)$$

$$\mathbb{C}_N[X] = \mathbb{C}[X]/(X^N + 1)$$

Examples (Real): $N = 2$

$$(1.2 + 2.3X) \cdot (3.2 + 4.1X) = 3.84 + 12.28X + 9.43X^2 = 12.28X - 5.59 \bmod (X^2 + 1)$$

$(\mathbb{Z}_N[X], +, \times)$, $(\mathbb{R}_N[X], +, \times)$ and $(\mathbb{C}_N[X], +, \times)$ are well defined as rings

- ✓ $(\mathbb{Z}_N[X], +)$, $(\mathbb{R}_N[X], +)$ and $(\mathbb{C}_N[X], +)$ are groups
- ✓ Multiplication $x \times y$ is well-defined!

Torus \mathbb{T} and Torus Polynomials $\mathbb{T}_N[X]$

$$\mathbb{T} = \mathbb{R}/\mathbb{Z}$$

$(\mathbb{T}, +, \cdot)$ is a \mathbb{Z} -module ($\cdot : \mathbb{Z} \times \mathbb{T} \rightarrow \mathbb{T}$ a valid external product)

- ✓ It is a group $x + y \pmod{\mathbb{Z}}$, and $-x \pmod{\mathbb{Z}}$
- ✓ It is a \mathbb{Z} -module: $3 \cdot 0.6 = 0.8 \pmod{\mathbb{Z}}$ is defined!
- ✗ It is **not** a ring: 0×0.6 is **not** defined!

$\mathbb{T}_N[X] = \mathbb{R}[X]/(X^N + 1) \pmod{\mathbb{Z}}$: polynomials with coeffs $\in \mathbb{R}/\mathbb{Z} \pmod{X^N + 1}$

$(\mathbb{T}_N[X], +, \cdot)$ is a $\mathbb{Z}_N[X]$ -module

- $(2X + 3) \cdot (0.4X + 0.5) = (0.2X + 0.7) \pmod{X^2 + 1} \pmod{\mathbb{Z}}$
- external product by integers polynomial

Torus \mathbb{T} and Torus Polynomials $\mathbb{T}_N[X]$

$$\mathbb{T} = \mathbb{R}/\mathbb{Z}$$

$(\mathbb{T}, +, \cdot)$ is a \mathbb{Z} -module ($\cdot : \mathbb{Z} \times \mathbb{T} \rightarrow \mathbb{T}$ a valid external product)

- ✓ It is a group $x + y \pmod{\mathbb{Z}}$, and $-x \pmod{\mathbb{Z}}$
- ✓ It is a \mathbb{Z} -module: $3 \cdot 0.6 = 0.8 \pmod{\mathbb{Z}}$ is defined!
- ✗ It is **not** a ring: 0×0.6 is **not** defined!

$\mathbb{T}_N[X] = \mathbb{R}[X]/(X^N + 1) \pmod{\mathbb{Z}}$: polynomials with coeffs $\in \mathbb{R}/\mathbb{Z} \pmod{X^N + 1}$

$(\mathbb{T}_N[X], +, \cdot)$ is a $\mathbb{Z}_N[X]$ -module

- $(2X + 3) \cdot (0.4X + 0.5) = (0.2X + 0.7) \pmod{X^2 + 1} \pmod{\mathbb{Z}}$
- external product by integers polynomial

TFHE Scheme

Consists of three encryption schemes:

- TLWE ciphertext: $\mu \in \mathbb{T} \mapsto (a, b := \mu + \langle a, s \rangle + e)$, $a \in_R \mathbb{T}^n$, $s \in \{0, 1\}^n$
- TRLWE ciphertext: $\mu \in \mathbb{T}_N[X] \mapsto (a, b := \mu + s \cdot a + e)$, $a \in_R \mathbb{T}_N[X]^k$, $s \in \mathbb{B}_N[X]$
- TRGSW ciphertext: encrypts elements of $\mathbb{Z}_N[X]$ with small norm

	message	ciphertext	key	lin. combin.	product
TLWE	\mathbb{T}	\mathbb{T}^{n+1}	\mathbb{B}^n	✓	✗
TRLWE	$\mathbb{T}_N[X]$	$\mathbb{T}_N[X]^{k+1}$	$\mathbb{B}_N[X]^k$	✓	✗
TRGSW	$\mathbb{Z}_N[X]$	vector of TRLWE	$\mathbb{B}_N[X]^k$	✓	✓

⊕ Internal TRGSW Product : $\boxtimes: \text{TRGSW} \times \text{TRGSW} \rightarrow \text{TRGSW}$

⊖ External product : $\boxtimes: \text{TRGSW} \times \text{TRLWE} \rightarrow \text{TRLWE}$

$$(\mu_A, \mu_B) \mapsto \mu_A \cdot \mu_B$$

$$(e_A, e_B) \mapsto \|\mu_A\|_1 \cdot e_B + O(e_A)$$

If $\|\mu_A\|_1 = 1$ the noise propagation is linear!

TFHE Scheme

Consists of three encryption schemes:

- TLWE ciphertext: $\mu \in \mathbb{T} \mapsto (a, b := \mu + \langle a, s \rangle + e)$, $a \in_R \mathbb{T}^n$, $s \in \{0, 1\}^n$
- TRLWE ciphertext: $\mu \in \mathbb{T}_N[X] \mapsto (a, b := \mu + s \cdot a + e)$, $a \in_R \mathbb{T}_N[X]^k$, $s \in \mathbb{B}_N[X]$
- TRGSW ciphertext: encrypts elements of $\mathbb{Z}_N[X]$ with small norm

	message	ciphertext	key	lin. combin.	product
TLWE	\mathbb{T}	\mathbb{T}^{n+1}	\mathbb{B}^n	✓	✗
TRLWE	$\mathbb{T}_N[X]$	$\mathbb{T}_N[X]^{k+1}$	$\mathbb{B}_N[X]^k$	✓	✗
TRGSW	$\mathbb{Z}_N[X]$	vector of TRLWE	$\mathbb{B}_N[X]^k$	✓	✓

⊕ Internal TRGSW Product : $\boxtimes: \text{TRGSW} \times \text{TRGSW} \rightarrow \text{TRGSW}$

⊗ External product : $\boxtimes: \text{TRGSW} \times \text{TRLWE} \rightarrow \text{TRLWE}$

$$(\mu_A, \mu_B) \mapsto \mu_A \cdot \mu_B$$

$$(e_A, e_B) \mapsto \|\mu_A\|_1 \cdot e_B + O(e_A)$$

If $\|\mu_A\|_1 = 1$ the noise propagation is linear!

TFHE Scheme

Consists of three encryption schemes:

- TLWE ciphertext: $\mu \in \mathbb{T} \mapsto (a, b := \mu + \langle a, s \rangle + e)$, $a \in_R \mathbb{T}^n$, $s \in \{0, 1\}^n$
- TRLWE ciphertext: $\mu \in \mathbb{T}_N[X] \mapsto (a, b := \mu + s \cdot a + e)$, $a \in_R \mathbb{T}_N[X]^k$, $s \in \mathbb{B}_N[X]$
- TRGSW ciphertext: encrypts elements of $\mathbb{Z}_N[X]$ with small norm

	message	ciphertext	key	lin. combin.	product
TLWE	\mathbb{T}	\mathbb{T}^{n+1}	\mathbb{B}^n	✓	✗
TRLWE	$\mathbb{T}_N[X]$	$\mathbb{T}_N[X]^{k+1}$	$\mathbb{B}_N[X]^k$	✓	✗
TRGSW	$\mathbb{Z}_N[X]$	vector of TRLWE	$\mathbb{B}_N[X]^k$	✓	✓

① **Internal TRGSW Product** : \boxtimes : $\text{TRGSW} \times \text{TRGSW} \longrightarrow \text{TRGSW}$

② **External product** : \boxdot : $\text{TRGSW} \times \text{TRLWE} \longrightarrow \text{TRLWE}$

$$(\mu_A, \mu_B) \mapsto \mu_A \cdot \mu_B$$

$$(e_A, e_B) \mapsto \|\mu_A\|_1 \cdot e_B + O(e_A)$$

If $\|\mu_A\|_1 = 1$ the noise propagation is linear!

RLWE internal product ?

Internal product requires to evaluate a polynomial in s :

$$\mu_1 \boxtimes \mu_2 = (b_1 - sa_1)(b_2 - sa_2) = b_1b_2 - (b_1a_2 + b_2a_1)s + a_1a_2s^2.$$

The term s^2 :

- dedicated relinearization/keyswitch techniques (2011, ...)
- but in fact, TRGSW provides the multiplication by a secret s !

The meaning of b_1b_2, a_1a_2, \dots

- lift in $\mathbb{R}_N[X]$
- additional message space restrictions are required to make such product meaningful

TRLWE

small integer linear combinations
 $x + y, x - y$
 $a.x$ for public $a \in \mathbb{Z}_N[X]$

Homomorphic operations hierarchy

TRLWE

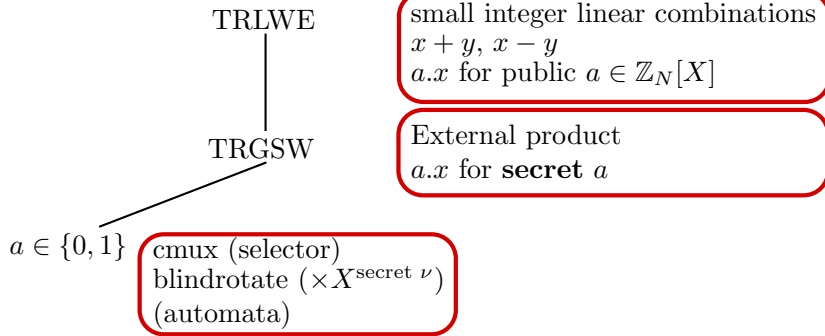


TRGSW

small integer linear combinations
 $x + y, x - y$
 $a.x$ for public $a \in \mathbb{Z}_N[X]$

External product
 $a.x$ for **secret** a

Homomorphic operations hierarchy



Homomorphic operations hierarchy

TRLWE

small integer linear combinations
 $x + y, x - y$
 $a.x$ for public $a \in \mathbb{Z}_N[X]$

TRGSW

External product
 $a.x$ for **secret** a

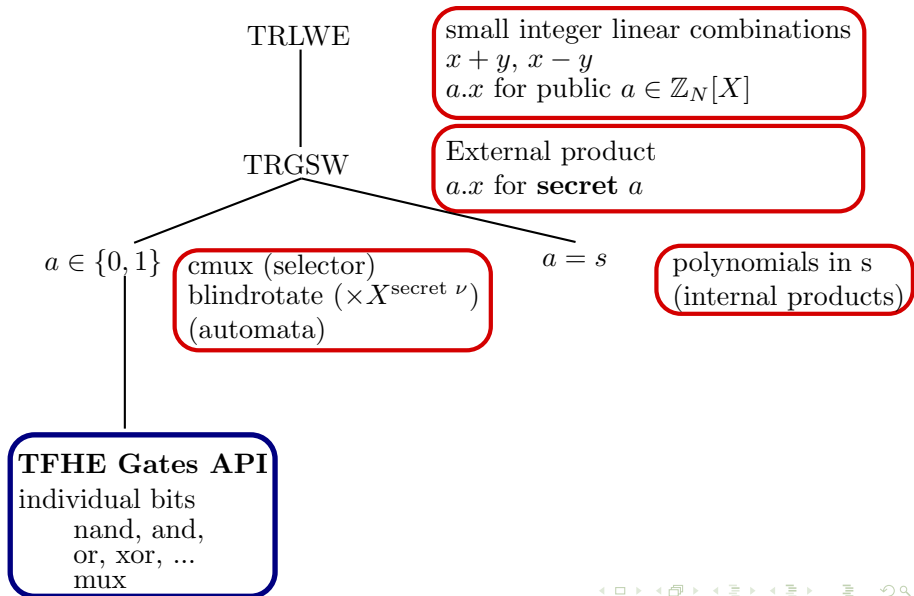
 $a \in \{0, 1\}$

cmux (selector)
 blindrotate ($\times X^{\text{secret } \nu}$)
 (automata)

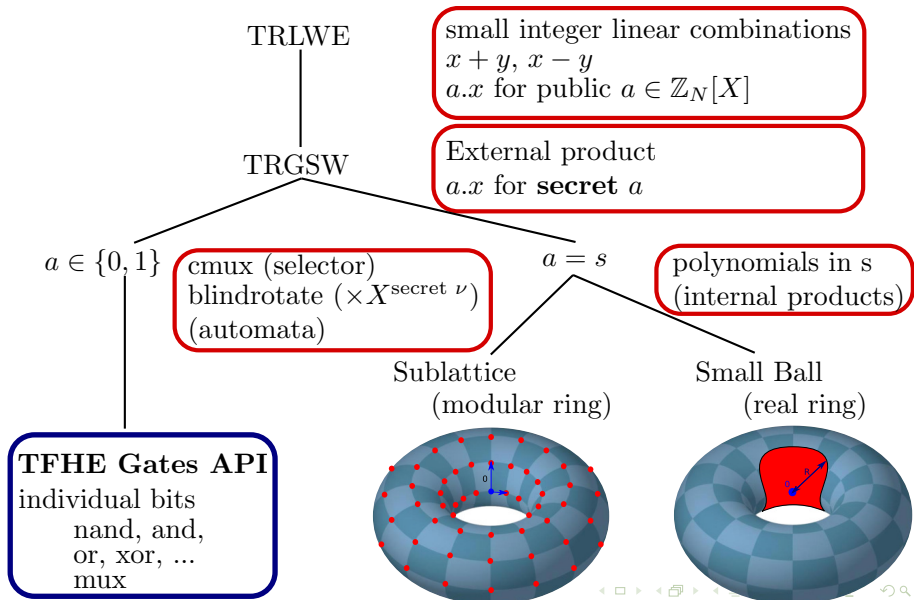
TFHE Gates API

individual bits
 nand, and,
 or, xor, ...
 mux

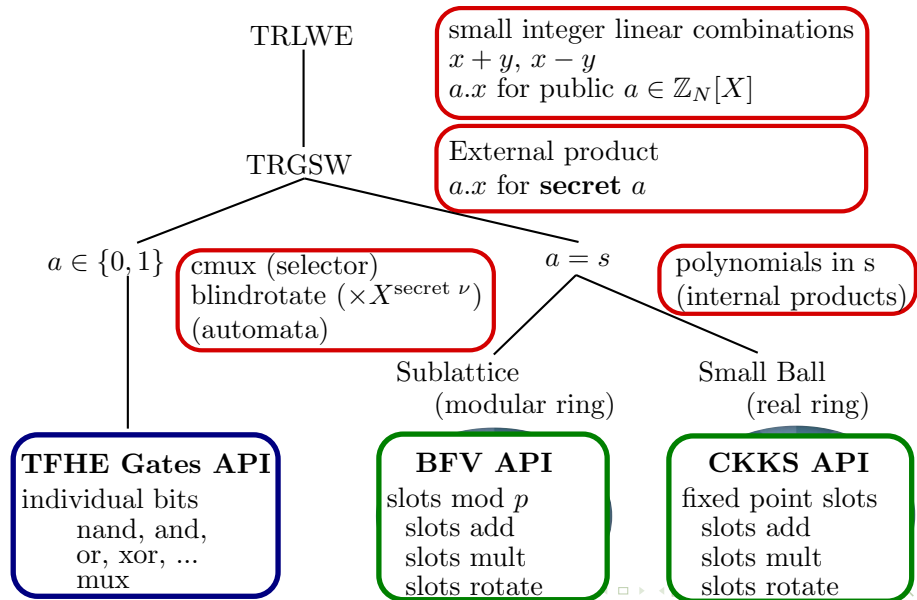
Homomorphic operations hierarchy



Homomorphic operations hierarchy




Homomorphic operations hierarchy



Different Models of Computations

- 1 TFHE: Binary circuit evaluation, LUTs, DFAs ...




- 2 B/FV: Integer arithmetic (SIMD)

decimal

$$\begin{array}{r}
 0011 \leftarrow \text{carries} \\
 4567 \\
 \underline{366} \\
 4933
 \end{array}$$

- 3 CKKS: Approximated (fixed-point) computations (SIMD)





Chimera: Combining different FHE schemes: TFHE, B/FV and CKKS

- Unified plaintext space over the Torus
- Switch between ciphertext representations (coefficient vs slot packing)

Coefficient and Slot packing

Coefficient packing

$$\mathbf{m} = \sum_{i=0}^{N-1} m_i \cdot X^i \quad \sim \quad \mathbf{m} = (m_0, m_1, \dots, m_{N-1})$$

with $m_i \in \mathbb{C}$ for all $i = 0, 1, \dots, N - 1$

m_0	m_1	m_2	\dots	m_{N-2}	m_{N-1}
-------	-------	-------	---------	-----------	-----------

Slot packing

$$X^N + 1 = \prod_{i=0}^{N-1} (X - \omega_i) \quad \sim \quad \mathbf{m} = (\mathbf{m}(\omega_0), \mathbf{m}(\omega_1), \dots, \mathbf{m}(\omega_{N-1}))$$

with $\omega_i \in \mathbb{C}$ for all $i = 0, 1, \dots, N - 1$

$\mathbf{m}(\omega_0)$	$\mathbf{m}(\omega_1)$	$\mathbf{m}(\omega_2)$	\dots	$\mathbf{m}(\omega_{N-2})$	$\mathbf{m}(\omega_{N-1})$
------------------------	------------------------	------------------------	---------	----------------------------	----------------------------

Coefficient and Slot packing

Coefficient packing

$$\mathbf{m} = \sum_{i=0}^{N-1} m_i \cdot X^i \quad \sim \quad \mathbf{m} = (m_0, m_1, \dots, m_{N-1})$$

with $m_i \in \mathbb{C}$ for all $i = 0, 1, \dots, N - 1$

m_0	m_1	m_2	\dots	m_{N-2}	m_{N-1}
-------	-------	-------	---------	-----------	-----------

Slot packing

$$X^N + 1 = \prod_{i=0}^{N-1} (X - \omega_i) \quad \sim \quad \mathbf{m} = (\mathbf{m}(\omega_0), \mathbf{m}(\omega_1), \dots, \mathbf{m}(\omega_{N-1}))$$

with $\omega_i \in \mathbb{C}$ for all $i = 0, 1, \dots, N - 1$

$\mathbf{m}(\omega_0)$	$\mathbf{m}(\omega_1)$	$\mathbf{m}(\omega_2)$	\dots	$\mathbf{m}(\omega_{N-2})$	$\mathbf{m}(\omega_{N-1})$
------------------------	------------------------	------------------------	---------	----------------------------	----------------------------


Morphism

There exists morphism to switch between the coefficient and slot representation!
(Vandermonde, DFT,...)

$$VDM = \begin{bmatrix} 1 & \omega_0^1 & \cdots & \omega_0^{N-1} \\ 1 & \omega_1^1 & \cdots & \omega_1^{N-1} \\ \vdots & \vdots & \cdots & \vdots \\ 1 & \omega_{N-1}^1 & \cdots & \omega_{N-1}^{N-1} \end{bmatrix}.$$

- A complex polynomial $\text{mod } X^N + 1$ carries N complex slots.
- A real polynomial $\text{mod } X^N + 1$ carries $N/2$ complex slots.
- Attention, some additional constraints are needed to define slots for $\mathbb{Z}_N[X]$.

How we can represent all plaintexts over the $\mathbb{T}_N[X]$? SANDBOXAQ inpher




BFV scheme (encoding)

- $\mathbb{Z}_N[X] \bmod p$: the ring of polynomials with integer $\bmod p$ coefficients module $X^N + 1$
- If $X^N + 1$ has N roots $\bmod p$, $\mathbb{Z}/p\mathbb{Z}^N$ is isomorphic to $\mathbb{Z}_N[X] \bmod p$

$$(\mathbb{Z}/p\mathbb{Z})^N \simeq \mathbb{Z}_N[X] \bmod p \simeq \frac{1}{p} \mathbb{Z}_N[X] \bmod \mathbb{Z}$$

The plaintext space \mathcal{M} is composed by exact multiples of $\frac{1}{p}$.



Plaintext addition $(\mu_1(X), \mu_2(X))$

$$\mu_1(X) + \mu_2(X) := \mu_1(X) + \mu_2(X) \bmod \mathbb{Z}.$$

Plaintext product (Montgomery) $(\mu_1(X), \mu_2(X))$

$$\mu_1(X) \boxtimes_p \mu_2(X) := p \cdot \tilde{\mu}_1(X) \cdot \tilde{\mu}_2(X) \bmod \mathbb{Z}, \text{ for lifts } \tilde{\mu}_1 \text{ and } \tilde{\mu}_2 \text{ in } \mathbb{R}_N[X]$$

Problem of lift

Examples: $p = 3$, $\mu_1 = \frac{1}{3}$ and $\mu_2 = \frac{2}{3}$

- Exact product: $3(I_1 + \frac{1}{3})(I_2 + \frac{2}{3}) = I + \frac{2}{3} = +\frac{2}{3} \pmod{1}$, for all I_1, I_2 integers
- Product with noise and small element: $3 * 5.33333 * 10.66665 = 170.6662$
- Product with noise and big element:
 $3 * 12345678.33333 * 7654321.66665 = -.839\dots$

- We need a small representative of the plaintext to keep the result correct.
- We should lift the ciphertext to small representative in $\mathbb{R}_N[X]$ (all coefficients in $[-1/2, 1/2)$).
- $\frac{1}{p} \gg \text{noise}$

Homomorphic operations

Homomorphic addition $c_1 = (a_1, b_1), c_2 = (a_2, b_2)$

$$(a, b) = (a_1 + a_2, b_1 + b_2)$$

Homomorphic product $c_1 = (a_1, b_1), c_2 = (a_2, b_2)$

$$\mu_1 = b_1 - s \cdot a_1 \text{ and } \mu_2 = b_2 - s \cdot a_2$$

$$\begin{aligned} \mu_1 \boxtimes_p \mu_2 &= p(\tilde{b}_1 - s \cdot \tilde{a}_1)(\tilde{b}_2 - s \cdot \tilde{a}_2) \\ &= \underbrace{(p \cdot \tilde{b}_1 \cdot \tilde{b}_2)}_{C_0} - s \cdot \underbrace{(p \cdot \tilde{a}_1 \cdot \tilde{b}_2 + p \cdot \tilde{a}_2 \cdot \tilde{b}_1)}_{C_1} + s^2 \cdot \underbrace{(p \cdot \tilde{a}_1 \cdot \tilde{a}_2)}_{C_2} \\ &= (b - s \cdot a) \end{aligned}$$


The term s^2 : relinearization with TRGSW encryption of $s!$

$$c_1 \boxtimes_p c_2 = (C_1, C_0) - TRGSW(s) \boxtimes (C_2, 0)$$

The meaning of $a_1 a_2, b_1 b_2 \dots$:

- small representatives in $\mathbb{R}_N[X]$

Fixed point



Fixed-point and Floating-points Numbers

Floating point (float, double in C):



- $x = m \cdot 2^\tau$, with $m \in 2^{-\rho} \cdot \mathbb{Z}$ and $\frac{1}{2} \leq |m| < 1$
- $\tau = \lceil \log_2(x) \rceil$ data dependent and **not public (not FHE-friendly)**
- **The exponent is always in sync with the data**
ex: $(1.23 \cdot 10^{-4}) * (7.24 \cdot 10^{-4}) = (8.90 \cdot 10^{-8})$

Fixed point:

- $x = m \cdot 2^\tau$, with $m \in 2^{-\rho} \cdot \mathbb{Z}$ and $0 \leq |m| < 1$,
- τ **is public, thus FHE-friendly**
- **Risk of overflow** (τ too small)
- **Risk of underflow** (τ too large)
ex: $(0.000123 \cdot 10^0) * (0.000724 \cdot 10^0) = (0.000000 \cdot 10^0)$

Addition is much trickier than you think!


- Given (m_1, τ_1) , (m_2, τ_2) , and τ .
- How do you compute $m \cdot 2^\tau = m_1 \cdot 2^{\tau_1} + m_2 \cdot 2^{\tau_2}$ with ρ bits of precision?
- Addition requires right shift and roundings, which are non-linear!



Continuous approach

- $x \times y = \text{Lift}(x) * \text{Lift}(y) \bmod \mathbb{Z}$.
- ✓ This approach can preserve (or reduce) the interval $[-\frac{1}{2L}, \frac{1}{2L}]$
- ✓ Lift is a periodic function: approx by sinus (or other Fourier serie) wherever it matters...
- ✗ ...but sinus can only be approx by a polynomial, which recursively requires a product.

Fixed point: CKKS



Discrete approach

- round a, b (and thus μ) on exact multiples of $\frac{1}{q}$ where $q \approx 2^{L+\rho}$.
- ✓ Brings us in the ring $\frac{1}{q}\mathbb{Z}_N[X] \bmod \mathbb{Z}$
- ✓ Exact Montgomery product $q(b_1 - sa_1)(b_2 - sa_2)$
 - The meaning of $a_1 a_2, b_1 b_2 \dots$: a_i, b_i are exact multiples of $\frac{1}{q}$
- ✗ Blows up the interval $[-\frac{1}{2L}, \frac{1}{2L}] \rightarrow [-\frac{1}{2^{L-\rho}}, \frac{1}{2^{L-\rho}}] \dots$
...works a leveled number of times.

Thank you for your attention!

Questions?

Appendix: Circuit Bootstrap mode versus Gate Bootstrap mode

TFHE in Circuit Bootstrap mode
Bootstrap after many gates
 (This work)

Input/Output

Plaintext (TLWE) → Ciphertext
 (TRGSW)
 Bit Overhead: 262144

- **Very fast** : transition in 34 μ s
- **No so fast**: circuit bootstrapped in 134 ms but after many gates
- **Composition** : LUT, (W)DFA

TFHE in Circuit bootstrap mode can evaluate LUT 16 to 8 in 1 sec

TFHE in Gate Bootstrap mode
Bootstrap between each gate
 (TFHE 2016 + optimizations)

Input/Output

Plaintext (TLWE) → Ciphertext
 (TLWE)
 Bit Overhead: 8000

- **No so fast**: bootstrapped binary gate runs in 13 ms
- All binary gates have the same cost
- **Composition**: unlimited

With TFHE we can compute 76 gates per second, for any circuit.

Appendix: Circuit Bootstrap mode versus Gate Bootstrap mode

TFHE in Circuit Bootstrap mode
Bootstrap after many gates
 (This work)

Input/Output

Plaintext (TLWE) \rightarrow Ciphertext
 (TRGSW)
 Bit Overhead: 262144

- **Very fast** : transition in $34 \mu\text{s}$
- **No so fast**: circuit bootstrapped in 134 ms but after many gates
- **Composition** : LUT, (W)DFA

TFHE in Circuit bootstrap mode can evaluate LUT 16 to 8 in 1 sec

TFHE in Gate Bootstrap mode
Bootstrap between each gate
 (TFHE 2016 + optimizations)

Input/Output

Plaintext (TLWE) \rightarrow Ciphertext
 (TLWE)
 Bit Overhead: 8000

- **No so fast**: bootstrapped binary gate runs in 13 ms
- All binary gates have the **same cost**
- **Composition**: unlimited

With TFHE we can compute 76 gates per second, for any circuit.